

Endogeneity, Instrumental Variables, Two-Stage Least Squares, and GMM

Eliab Luvanda

2026-04-26

Table of contents

1	Introduction and Motivation	4
1.1	The Classical Linear Model and Its Assumptions	4
1.2	Endogeneity: Definition, Sources, and Consequences	5
1.2.1	Omitted Variable Bias	5
1.2.2	Simultaneity Bias	6
1.2.3	Measurement Error	6
1.3	A Preview of Solutions	7
2	Instrumental Variable Estimator	7
3	Two-Stage Least Squares (2SLS) Estimator	10
3.1	Variance-covariance Matrix of the 2SLS Estimator	11
3.1.1	Variance-covariance Matrix under Spherical Errors	11
3.1.2	Variance-covariance Matrix under Heteroskedasticity	12
3.1.3	Variance-Covariance Matrix under Serial Correlation	12
3.2	Properties of the 2SLS Estimator	12
3.2.1	Consistency	12
3.2.2	Asymptotic Normality	13
3.2.3	Efficiency	14
3.3	Testing for Endogeneity	14
3.4	Weak Instruments	15
3.5	Overidentification and the Sargan-Hansen Test	15
4	Generalized Method of Moments (GMM) Estimator	15
4.1	The Weighting Matrix $\mathbf{W}_n$	17
4.2	The Optimal Weighting Matrix	17
4.3	Second Step GMM Estimation	18
4.4	Properties of the GMM Estimators	19
4.4.1	Consistency	19
4.4.2	Asymptotic Normality	22
4.4.3	Efficiency	23
4.5	Tests	24
4.5.1	Overidentifying restrictions	24
4.6	Other Diagnostic Tests	24
5	Synthesis and Summary	25

List of Figures	3
-----------------	---

6 Review Questions and Exercises	28
---	-----------

A Monte Carlo Simulation: OLS, IV, 2SLS, and GMM under Endogeneity	33
---	-----------

A.1 Simulation Results	33
----------------------------------	----

A.2 R Code	35
----------------------	----

A.3 Documentation of the R Code	40
---	----

List of Figures

1 Empirical densities of the Monte Carlo distribution of $\hat{\beta}_2$ for OLS, simple IV, 2SLS, and GMM.	34
---	----

List of Tables

1 Summary of the simple IV, 2SLS, and GMM estimators	27
2 Mean, variance, and bias of the Monte Carlo distribution of $\hat{\beta}_2$	35

1. Introduction and Motivation

The empirical study of economic relationships rests, almost without exception, on some version of the regression model. Yet the conditions under which regression yields reliable estimates of structural parameters are considerably more stringent than is often acknowledged in applied work. This section lays out the classical framework, identifies the precise point at which it breaks down, and sketches the remedies that the remainder of these notes develop in full; standard treatments include Hayashi [11], Wooldridge [16], and Hansen [10].

1.1 The Classical Linear Model and Its Assumptions

Let $\{(y_i, \mathbf{x}_i)\}_{i=1}^n$ be a sample of n observations, where $y_i \in \mathbb{R}$ is the dependent variable and $\mathbf{x}_i \in \mathbb{R}^k$ is a vector of regressors, with the convention that its first component equals unity. The *classical linear model* (CLM) postulates

$$y_i = \mathbf{x}_i^\top \boldsymbol{\beta}_0 + \varepsilon_i, \quad i = 1, \dots, n, \quad (1)$$

where $\boldsymbol{\beta}_0 \in \mathbb{R}^k$ is the vector of true parameters and ε_i is an unobservable disturbance. Written in matrix form, (1) becomes $\mathbf{Y} = \mathbf{X}\boldsymbol{\beta}_0 + \boldsymbol{\varepsilon}$, where $\mathbf{Y} \in \mathbb{R}^n$, $\mathbf{X} \in \mathbb{R}^{n \times k}$, and $\boldsymbol{\varepsilon} \in \mathbb{R}^n$ collect the observations in the obvious way.

The ordinary least-squares (OLS) estimator is

$$\hat{\boldsymbol{\beta}}_{\text{OLS}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Y} = \boldsymbol{\beta}_0 + (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\boldsymbol{\varepsilon}. \quad (2)$$

Its desirable properties depend on the following assumptions, which we state in a form suitable for asymptotic analysis.

Assumption 1.1 (Linearity). The data-generating process satisfies (1) for some $\boldsymbol{\beta}_0 \in \mathbb{R}^k$.

Assumption 1.2 (Random Sampling). $\{(y_i, \mathbf{x}_i)\}_{i=1}^n$ is an independent and identically distributed (i.i.d.) sequence drawn from a well-defined population distribution.

Assumption 1.3 (Full Rank). $\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top]$ is finite and positive definite, so that $\text{rank}(\mathbf{X}'\mathbf{X}) = k$ almost surely for n sufficiently large.

Assumption 1.4 (Strict Exogeneity). $\mathbb{E}[\varepsilon_i \mid \mathbf{x}_i] = 0$ almost surely. As an immediate consequence, $\mathbb{E}[\mathbf{x}_i \varepsilon_i] = \mathbf{0}$ and $\mathbb{E}[\varepsilon_i] = 0$.

Assumption 1.5 (Spherical Disturbances). $\mathbb{E}[\varepsilon_i^2 \mid \mathbf{x}_i] = \sigma^2 \in (0, \infty)$ and $\mathbb{E}[\varepsilon_i \varepsilon_j \mid \mathbf{x}_i, \mathbf{x}_j] = 0$ for all $i \neq j$.

Remark 1.6. Assumption 1.4 is the pivot around which the entire analysis turns. As stated, it is a *mean independence* condition: $\mathbb{E}[\varepsilon_i \mid \mathbf{x}_i] = 0$. This condition implies the weaker orthogonality condition $\mathbb{E}[\mathbf{x}_i \varepsilon_i] = \mathbf{0}$, but the converse does not hold in general. For consistency of OLS, the orthogonality condition is the key requirement. Stronger conditional mean assumptions are mainly useful for deriving sharper finite-sample or conditional results.

Under Assumptions~1.1–1.5, the Gauss–Markov theorem establishes that OLS is the best linear unbiased estimator (BLUE) of β_0 . Consistency and asymptotic normality follow from a standard application of the law of large numbers and the central limit theorem; see Wooldridge [16], Hayashi [11], and Hansen [10].

Theorem 1.7 (Consistency and Asymptotic Normality of OLS). *Under Assumptions 1.1–1.4 and the moment condition $\mathbb{E}[\|\mathbf{x}_i\|^2 \varepsilon_i^2] < \infty$,*

- (i) $\hat{\beta}_{\text{OLS}} \xrightarrow{p} \beta_0$; and
- (ii) $\sqrt{n}(\hat{\beta}_{\text{OLS}} - \beta_0) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_{\text{OLS}})$, where $\Sigma_{\text{OLS}} = (\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top])^{-1} \mathbb{E}[\varepsilon_i^2 \mathbf{x}_i \mathbf{x}_i^\top] (\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top])^{-1}$.

Under the additional Assumption 1.5, Σ_{OLS} simplifies to $\sigma^2 (\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top])^{-1}$.

Proof. From (2),

$$\hat{\beta}_{\text{OLS}} - \beta_0 = \left(\frac{1}{n} \mathbf{X}' \mathbf{X} \right)^{-1} \frac{1}{n} \mathbf{X}' \varepsilon.$$

By the LLN (Assumption 1.2), $n^{-1} \mathbf{X}' \mathbf{X} \xrightarrow{p} \mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top]$, which is invertible by Assumption 1.3, and $n^{-1} \mathbf{X}' \varepsilon \xrightarrow{p} \mathbb{E}[\mathbf{x}_i \varepsilon_i] = \mathbf{0}$ by Assumption 1.4. Part (i) follows from Slutsky's theorem. For part (ii), write $n^{-1/2} \mathbf{X}' \varepsilon = n^{-1/2} \sum_{i=1}^n \mathbf{x}_i \varepsilon_i$. The summands are i.i.d. with mean zero and covariance matrix $\mathbb{E}[\varepsilon_i^2 \mathbf{x}_i \mathbf{x}_i^\top]$; the CLT yields $n^{-1/2} \mathbf{X}' \varepsilon \xrightarrow{d} \mathcal{N}(\mathbf{0}, \mathbb{E}[\varepsilon_i^2 \mathbf{x}_i \mathbf{x}_i^\top])$. The delta method (or Slutsky) completes the argument. \$■\$

1.2 Endogeneity: Definition, Sources, and Consequences

The entire edifice of Theorem~1.7 collapses once Assumption~1.4 is violated.

Definition 1.8 (Endogeneity). A regressor x_{ji} (the j -th component of $\mathbf{x}_i$) is said to be *endogenous* if $\text{Cov}(x_{ji}, \varepsilon_i) \neq 0$, or, more generally, if $\mathbb{E}[\mathbf{x}_i \varepsilon_i] \neq \mathbf{0}$. The model (1) is then said to suffer from *endogeneity*; compare Wooldridge [16], Cameron and Trivedi [2], and Hansen [10].

Endogeneity is not a single pathology but a family of related problems sharing a common algebraic signature. We distinguish three principal sources.

1.2.1 Omitted Variable Bias

Suppose the true model is

$$y_i = \mathbf{x}_i^\top \beta_0 + \gamma_0 c_i + u_i, \quad \mathbb{E}[u_i | \mathbf{x}_i, c_i] = 0,$$

where c_i is an unobserved confounder correlated with $\mathbf{x}_i$. If we estimate the misspecified model (1) with $\varepsilon_i = \gamma_0 c_i + u_i$, then $\mathbb{E}[\mathbf{x}_i \varepsilon_i] = \gamma_0 \mathbb{E}[\mathbf{x}_i c_i] \neq \mathbf{0}$ whenever $\gamma_0 \neq 0$ and c_i is correlated with at least one regressor. This is the *omitted variable bias* (OVB) problem, arguably the most common form of endogeneity in applied economics; see Wooldridge [16] and Cameron and Trivedi [2].

1.2.2 Simultaneity Bias

In many economic settings, y_i and one or more components of $\mathbf{x}_i$ are determined jointly by a system of equations. Hayashi [11] and Davidson and MacKinnon [7] discuss the canonical supply-and-demand example:

$$\begin{aligned} q_i^d &= \alpha_0 + \alpha_1 p_i + \alpha_2 \mathbf{w}_i' \boldsymbol{\delta} + u_i^d \quad (\text{demand}), \\ q_i^s &= \beta_0 + \beta_1 p_i + \beta_2 \mathbf{r}_i' \boldsymbol{\phi} + u_i^s \quad (\text{supply}), \\ q_i^d &= q_i^s \quad (\text{market clearing}). \end{aligned}$$

In equilibrium, price p_i is a function of both structural errors u_i^d and u_i^s . Regressing quantity on price therefore yields an estimator that is inconsistent for either structural parameter.

1.2.3 Measurement Error

Let x_i^* be the true value of a regressor, observed with error: $x_i = x_i^* + v_i$, where v_i is classical measurement error with $\mathbb{E}[v_i] = 0$, $\text{Cov}(x_i^*, v_i) = 0$. As Wooldridge [16] and Cameron and Trivedi [2] show, substituting into the true model $y_i = \beta_0 + \beta_1 x_i^* + u_i$ gives

$$y_i = \beta_0 + \beta_1 x_i + (u_i - \beta_1 v_i). \\ \text{[REDACTED]} \{ \text{[REDACTED]} \} \\ = \varepsilon_i$$

Since $\text{Cov}(x_i, \varepsilon_i) = -\beta_1 \text{Var}(v_i) \neq 0$, OLS is inconsistent. The resulting *attenuation bias* pushes the estimate of β_1 toward zero in absolute value.

The consequences of endogeneity are stark.

Theorem 1.9 (Inconsistency of OLS under Endogeneity). *Suppose Assumption 1.4 fails, so that $\boldsymbol{\delta} \equiv \mathbb{E}[\mathbf{x}_i \varepsilon_i] \neq \mathbf{0}$. Then, under Assumptions 1.1–1.3,*

$$\text{plim}_{n \rightarrow \infty} \hat{\boldsymbol{\beta}}_{\text{OLS}} = \boldsymbol{\beta}_0 + (\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top])^{-1} \boldsymbol{\delta} \neq \boldsymbol{\beta}_0.$$

Moreover, no amount of additional data can eliminate this asymptotic bias.

Proof. By the LLN applied to (2),

$$\text{plim } \hat{\boldsymbol{\beta}}_{\text{OLS}} = \boldsymbol{\beta}_0 + \left(\text{plim } \frac{\mathbf{X}'\mathbf{X}}{n} \right)^{-1} \left(\text{plim } \frac{\mathbf{X}'\boldsymbol{\varepsilon}}{n} \right) = \boldsymbol{\beta}_0 + (\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top])^{-1} \boldsymbol{\delta}.$$

Because $\boldsymbol{\delta} \neq \mathbf{0}$ by hypothesis and $\mathbb{E}[\mathbf{x}_i \mathbf{x}_i^\top]$ is non-singular by Assumption 1.3, the second term is non-zero and persists as $n \rightarrow \infty$. \$■\$

Remark 1.10. Theorem 1.9 makes clear that endogeneity is fundamentally an identification problem, not a finite-sample one. Collecting more data merely produces a more precise estimate of the wrong quantity. The solution must address the source of the correlation $\boldsymbol{\delta}$, not the sample size.

Example 1.11 (Returns to Schooling). Consider the classic Mincer earnings equation,

$$\log w_i = \beta_0 + \beta_1 S_i + \mathbf{x}_i^\top \boldsymbol{\gamma} + \varepsilon_i,$$

where w_i is the wage, S_i is years of schooling, and $\mathbf{x}_i$ contains observable controls (age, region, etc.). The parameter of interest, β_1 , measures the average causal return to an additional year of education. The problem is that S_i is almost certainly correlated with unobserved *ability* a_i , which also affects wages and is absorbed into ε_i . OLS therefore overstates β_1 . Card [3] addresses this by instrumenting S_i with proximity to a college, a variable that plausibly affects schooling attainment without directly influencing wages — an application of the IV strategy introduced in Section 2.

1.3 A Preview of Solutions

Three broad strategies have been developed to recover consistent estimates of β_0 in the presence of endogeneity. These notes treat the first two in depth and briefly touch on their connections to the third.

Instrumental Variables (IV). The IV approach introduces a vector of *instruments* $\mathbf{z}_i \in \mathbb{R}^\ell$ that satisfy two conditions: *relevance*, $\text{Cov}(\mathbf{z}_i, \mathbf{x}_i) \neq \mathbf{0}$, and *exogeneity*, $\mathbb{E}[\mathbf{z}_i \varepsilon_i] = \mathbf{0}$. As Hayashi [11], Hansen [10], and Wooldridge [16] explain, these two conditions together restore identification. Section~2 develops the IV estimator, establishes its asymptotic properties, and discusses the critical practical issue of *instrument strength*.

Two-Stage Least Squares (2SLS). When $\ell > k$ — i.e. when there are more instruments than endogenous regressors — the model is *overidentified* and IV is no longer unique. Amemiya [1], Hayashi [11], and Wooldridge [16] show that 2SLS provides an optimal way to exploit all available instruments simultaneously. Section~3 derives the 2SLS estimator, relates it to IV in the just-identified case, and establishes its efficiency within the class of IV estimators under homoskedastic errors.

Generalized Method of Moments (GMM). GMM nests both IV and 2SLS as special cases. It begins from a set of moment conditions $\mathbb{E}[\mathbf{g}(\mathbf{w}_i, \boldsymbol{\beta}_0)] = \mathbf{0}$ and constructs the estimator by minimising a quadratic form in the sample analogue. With an *optimal* weighting matrix, GMM is asymptotically efficient within the class of GMM estimators based on the same moment conditions. As Hayashi [11], Hall [8], and Hansen [10] emphasize, GMM also extends naturally to heteroskedastic and autocorrelated settings by changing the weighting matrix and variance estimation rather than the underlying moment conditions. Section~4 provides a self-contained treatment.

Diagnostics. Identifying endogeneity in a given dataset, assessing instrument validity, and testing overidentifying restrictions are tasks of independent importance. Section~3.3 covers the Hausman test, the Sargan–Hansen J -statistic, and the Stock–Yogo weak-instrument test.

2. Instrumental Variable Estimator

To address the endogeneity problem, we can introduce instrumental variables (IVs).

Definition 2.1 (Instrumental Variable). A variable z_i is an *instrument* for the endogenous regressor x_i if it satisfies the following conditions:

- (i) **Relevance:** $\text{Cov}(z_i, x_i) \neq 0$.
- (ii) **Exogeneity:** $\mathbb{E}[z_i \varepsilon_i] = 0$.

The *relevance condition* ensures that the instrument is correlated with the endogenous regressor. This guarantees that the instrument contains identifying information about the endogenous variation in x_i and that the IV estimator is well-defined in the just-identified case. As Hayashi [11], Hansen [10], and Wooldridge [16] note, stronger relevance typically yields a more precise estimator, whereas weak relevance can lead to poor finite-sample performance and unreliable inference.

The *exogeneity condition* ensures that the instrument is not correlated with the error term, which is crucial for the consistency of the IV estimator. As Hayashi [11], Hansen [10], and Wooldridge [16] make clear, if this condition is violated, the IV estimator will be biased and inconsistent.

Before deriving the IV estimator, let us consider the following linear regression model with one endogenous regressor and two exogenous regressors:

$$y_i = \beta_1 + \beta_2 x_2 + \mathbf{x}_i^\top \boldsymbol{\gamma} + \varepsilon_i, \quad (3)$$

where x_2 is the endogenous regressor, $\mathbf{x}_i$ is a vector of exogenous regressors, and ε_i is the error term. We assume that we have an instrument z_2 for x_2 that satisfies the relevance and exogeneity conditions. In this specification, β_1 is the intercept, β_2 is the coefficient on the endogenous regressor x_2 , and $\boldsymbol{\gamma}$ is a vector of coefficients on the exogenous regressors.

We can also write the model in matrix form as follows:

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad (4)$$

where $\mathbf{y}$ is the $n \times 1$ vector of the dependent variable, $\mathbf{X}$ is the $n \times k$ matrix of regressors (including the intercept, the endogenous regressor, and the exogenous regressors), $\boldsymbol{\beta}$ is the $k \times 1$ vector of parameters, and $\boldsymbol{\varepsilon}$ is the $n \times 1$ vector of error terms.

We can partition the matrix of regressors $\mathbf{X}$ into two parts: $\mathbf{X} = [\mathbf{X}_1 \mid \mathbf{X}_2]$, where $\mathbf{X}_1$ contains the exogenous regressors (including the intercept) and $\mathbf{X}_2$ contains the endogenous regressor x_2 . We can define the vector of instruments $\mathbf{Z} = [\mathbf{X}_1 \mid \mathbf{z}_2]$, where $\mathbf{z}_2$ is the $n \times 1$ vector of the instrument for x_2 . Since x_2 is endogenous, $E[\mathbf{X}'\boldsymbol{\varepsilon}] \neq \mathbf{0}$, because $E[x_2 \varepsilon_i] \neq 0$. However, since $\mathbf{Z}$ contains only exogenous variables, we have $E[\mathbf{Z}'\boldsymbol{\varepsilon}] = \mathbf{0}$. The IV estimator can be derived by exploiting the moment condition $E[\mathbf{Z}'\boldsymbol{\varepsilon}] = \mathbf{0}$. Specifically, the IV estimator $\hat{\boldsymbol{\beta}}_{\text{IV}}$ is defined as the solution to the following equation:

$$E[\mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}})] = \mathbf{0}. \quad (5)$$

We have substituted $\boldsymbol{\varepsilon}$ with $\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}}$ in the moment condition.

Using the principle of the method of moments estimation, which states that the sample analogue of the population moment condition should hold at the estimated parameters, we can proceed with the derivation of the IV estimator, as follows:

$$\begin{aligned}
1/n[\mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}})] &= \mathbf{0} \\
\mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}}) &= \mathbf{0} \\
\mathbf{Z}'\mathbf{y} - \mathbf{Z}'\mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}} &= \mathbf{0} \\
\mathbf{Z}'\mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}} &= \mathbf{Z}'\mathbf{y} \\
\hat{\boldsymbol{\beta}}_{\text{IV}} &= (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\mathbf{y}.
\end{aligned}$$

In the second line, we have multiplied both sides of the equation by n to eliminate the $1/n$ term. In the third line, we have rearranged the equation to isolate $\mathbf{Z}'\mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}}$ on one side. Finally, in the last line, we have solved for $\hat{\boldsymbol{\beta}}_{\text{IV}}$ by pre-multiplying both sides of the equation by $(\mathbf{Z}'\mathbf{X})^{-1}$.

Unlike OLS, which uses the orthogonality condition between regressors and errors, the IV estimator replaces that invalid condition with valid moment conditions involving the instruments $\mathbf{Z}$. In this sense, IV corrects the endogeneity problem by projecting identification onto exogenous variation in the instruments. The IV estimator is consistent for $\boldsymbol{\beta}$ when the instruments are valid, the model is correctly specified, and the relevant rank condition holds.

Proposition 2.2 (Consistency of the IV Estimator). *Under the assumptions of the linear model, random sampling, full rank, and the validity of the instruments (i.e., relevance and exogeneity), the IV estimator $\hat{\boldsymbol{\beta}}_{\text{IV}}$ is consistent for $\boldsymbol{\beta}$.*

Proof. We can write the IV estimator as follows:

$$\begin{aligned}
\hat{\boldsymbol{\beta}}_{\text{IV}} &= (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\mathbf{y} \\
&= (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'(\mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}) \\
&= \boldsymbol{\beta} + (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\boldsymbol{\varepsilon}.
\end{aligned}$$

Since $n^{-1}\mathbf{Z}'\mathbf{X} \xrightarrow{p} \mathbb{E}[\mathbf{z}_i\mathbf{x}_i']$, and the limit is non-singular by the relevance condition, the inverse exists with probability approaching one. Moreover, since $\mathbb{E}[\mathbf{z}_i\boldsymbol{\varepsilon}_i] = \mathbf{0}$ by the exogeneity condition, we have $n^{-1}\mathbf{Z}'\boldsymbol{\varepsilon} \xrightarrow{p} \mathbf{0}$. Therefore, by Slutsky's theorem, $\hat{\boldsymbol{\beta}}_{\text{IV}} \xrightarrow{p} \boldsymbol{\beta}$. \$■\$

Remark 2.3. The consistency of the IV estimator relies crucially on the validity of the instruments. If the instruments are only weakly correlated with the endogenous regressor, that is, if the relevance condition is not satisfied, then the IV estimator will be less precise and may even be inconsistent. If the instruments are correlated with the error term, that is, if the exogeneity condition is not satisfied, then the IV estimator will be biased and inconsistent. Therefore, it is essential to carefully select and validate instruments in applied work.

Remark 2.4. Normally, we do not prove the unbiasedness of the IV estimator, because it is not unbiased in small samples. However, it is consistent, which means that as the sample size increases, the bias of the IV estimator diminishes and it converges in probability to the true parameter value β .

The IV idea extends naturally to models with multiple endogenous regressors and multiple instruments. In the just-identified multivariate case, where the number of instruments equals the number of endogenous regressors, the estimator still takes the form $\hat{\beta}_{IV} = (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\mathbf{y}$. In the overidentified case, however, there are more instruments than endogenous regressors, and estimation is typically carried out using 2SLS or GMM rather than this simple just-identified formula. These issues are discussed in more detail in the subsequent sections on Two-Stage Least Squares (2SLS) and Generalized Method of Moments (GMM).

Remark 2.5. In the just-identified case, where the number of instruments equals the number of endogenous regressors, the IV estimator is equivalent to the Two-Stage Least Squares (2SLS) estimator. However, in the overidentified case, where there are more instruments than endogenous regressors, the 2SLS estimator provides a way to combine all the available instruments to achieve greater efficiency.

3. Two-Stage Least Squares (2SLS) Estimator

The two-stage least squares (2SLS) estimator is a generalization of the IV estimator that allows for multiple endogenous regressors and multiple instruments. Wooldridge [16], Hayashi [11], and Amemiya [1] provide standard treatments of the 2SLS estimator, including its derivation, properties, and applications. The 2SLS estimator is particularly useful in the overidentified case, where there are more instruments than endogenous regressors, as it allows for the efficient use of all available instruments. Following these treatments, we can specify the framework for deriving the 2SLS estimator as follows:

Consider a general linear regression model with k regressors. This model can be presented as follows:

$$y_i = \beta_1 + \beta_2 x_{2i} + \dots + \beta_k x_{ki} + \varepsilon_i, \quad (6)$$

where y_i is the dependent variable for observation i and ε_i is the corresponding error term. The model includes k regressors, including the intercept. Let us assume that $x_2, \dots, x_{k-1}$ are exogenous regressors, while x_k is an endogenous regressor. In this case, we have one endogenous regressor and $k - 1$ exogenous regressors. This implies that $E[x_k \varepsilon_i] \neq 0$, while $E[x_j \varepsilon_i] = 0$ for $j = 2, \dots, k - 1$.

The m instruments for the endogenous regressor x_k are denoted by $z_{1i}, \dots, z_{mi}$, and they satisfy the relevance condition $\text{Cov}(z_{ji}, x_k) \neq 0$ for $j = 1, \dots, m$, and the exogeneity condition $E[z_{ji} \varepsilon_i] = 0$ for $j = 1, \dots, m$. The matrix of instruments can be denoted as $\mathbf{Z} = [\mathbf{X}_1 \mid \mathbf{z}_k]$, where $\mathbf{X}_1$ contains the exogenous regressors (including the intercept) and $\mathbf{z}_k$ is the vector of instruments for the endogenous regressor x_k . The 2SLS estimator can be derived by first regressing the endogenous regressor x_k on the instruments $\mathbf{Z}$. This is the first stage of the 2SLS procedure, and it can be written as follows:

$$x_k = \mathbf{Z}\gamma + v, \quad (7)$$

where γ is a vector of parameters to be estimated, and v is the error term in the first stage regression.

In equation (7), x_k is regressed on the instruments $\mathbf{Z}$, which include the exogenous regressors and the instruments for the endogenous regressor. Thus, x_k is expressed as a linear combination of the instruments. If the parameter vector γ were known, and we ignored the error term v , we could recover the exogenous component of x_k as $\hat{x}_k = \mathbf{Z}\gamma$. However, since γ is unknown, we can consistently estimate it using OLS because the instruments are exogenous, that is, $E[v'\mathbf{Z}] = 0$. The OLS estimator for γ is given by $\hat{\gamma} = (\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'x_k$. We can then use this estimate to generate the fitted values of x_k , which are denoted by $\hat{x}_k = \mathbf{Z}\hat{\gamma}$. Substituting $(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'x_k$ for $\hat{\gamma}$ gives us $\hat{x}_k = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'x_k$. This is the predicted value of the endogenous regressor x_k based on the instruments $\mathbf{Z}$. Using the expression for $\hat{x}_k$, we can write the second-stage OLS estimator for β as follows:

$$\hat{\beta}_{2SLS} = (\hat{\mathbf{X}}'\hat{\mathbf{X}})^{-1}\hat{\mathbf{X}}'\mathbf{y}, \quad (8)$$

Substitution of the expression for $\hat{\mathbf{X}}$ into the second stage regression, together with some algebraic manipulation, yields the 2SLS estimator in the following form:

$$\begin{aligned} \hat{\beta}_{2SLS} &= (\hat{\mathbf{X}}'\hat{\mathbf{X}})^{-1}\hat{\mathbf{X}}'\mathbf{y} \\ &= (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y} \\ &= (\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{y}. \end{aligned}$$

where $\mathbf{P}_Z$ is the projection matrix onto the space spanned by $z_1, \dots, z_m$, defined as $\mathbf{P}_Z = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'$. As Wooldridge [16], Hayashi [11], and Amemiya [1] explain, the 2SLS estimator can be interpreted as the OLS estimator applied to the transformed data, where the endogenous regressor is replaced by its fitted values from the first stage regression. Under the same conditions used for IV consistency, 2SLS is consistent for β , and with multiple instruments it uses the available moment information more efficiently.

3.1 Variance-covariance Matrix of the 2SLS Estimator

3.1.1 Variance-covariance Matrix under Spherical Errors

As Hayashi [11], Amemiya [1], and Wooldridge [16] show, under homoskedasticity the finite-sample conditional variance of the 2SLS estimator has the familiar projection-matrix form

$$\text{Var}(\hat{\beta}_{2SLS} | \mathbf{X}, \mathbf{Z}) = \sigma^2(\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1},$$

where $\mathbf{P}_Z = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'$ and $\sigma^2 = \text{Var}(\varepsilon_i | \mathbf{Z})$ under the homoskedastic model. Equivalently, since $\hat{\beta}_{2SLS} = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y}$, inference under spherical errors is based on the matrix $(\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}$.

Here, σ^2 can be estimated using the residuals from the second stage regression as follows:

$$\hat{\sigma}^2 = \frac{1}{n-k} \sum_{i=1}^n (y_i - \hat{y}_i)^2,$$

where $\hat{y}_i$ is the fitted value of y_i from the second stage regression, and k is the number of parameters estimated in the second stage regression (including the intercept). The standard errors of the 2SLS estimator are then obtained from the square roots of the diagonal elements of $\hat{\sigma}^2(\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}$.

3.1.2 Variance-covariance Matrix under Heteroskedasticity

In the presence of heteroskedasticity, the homoskedastic formula above is no longer valid. A standard heteroskedasticity-robust covariance estimator, discussed by White [15] and Hansen [10], is the sandwich form

$$\widehat{\text{Var}}_{\text{rob}}(\hat{\beta}_{2\text{SLS}}) = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\hat{\mathbf{\Omega}}\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X}(\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1},$$

where $\hat{\mathbf{\Omega}}$ is an estimator of the covariance matrix of the structural disturbances. In the simplest heteroskedastic but uncorrelated case, one uses $\hat{\mathbf{\Omega}} = \text{diag}(\hat{\varepsilon}_1^2, \dots, \hat{\varepsilon}_n^2)$, where $\hat{\varepsilon}_i$ is the second-stage residual for observation i .

3.1.3 Variance-Covariance Matrix under Serial Correlation

With serial correlation, one replaces $\hat{\mathbf{\Omega}}$ in the robust sandwich formula by a heteroskedasticity-and-autocorrelation-consistent (HAC) estimator, such as the Newey–West estimator; see Hayashi [11] and Hansen [10]. The resulting covariance estimator is

$$\widehat{\text{Var}}_{\text{NW}}(\hat{\beta}_{2\text{SLS}}) = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\hat{\mathbf{\Omega}}_{\text{NW}}\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X}(\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1},$$

where $\hat{\mathbf{\Omega}}_{\text{NW}}$ denotes a HAC estimator that incorporates both heteroskedasticity and autocorrelation up to a chosen lag length.

3.2 Properties of the 2SLS Estimator

3.2.1 Consistency

Proposition 3.1 (Consistency of the 2SLS Estimator). *Under the assumptions of the linear model, random sampling, full rank, and the validity of the instruments, the 2SLS estimator $\hat{\beta}_{2\text{SLS}}$ is consistent for β .*

Proof. The consistency of the 2SLS estimator can be established by showing that it converges in probability to β as the sample size increases. We can write the 2SLS estimator as follows:

$$\begin{aligned}\hat{\beta}_{2\text{SLS}} &= (\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{y} \\ &= \beta + (\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\varepsilon.\end{aligned}$$

Since β is a constant vector, we can write its convergence to zero as follows:

$$\hat{\beta}_{2\text{SLS}} - \beta = (\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'\varepsilon.$$

We can show that $\text{plim } \hat{\beta}_{2\text{SLS}} - \beta = 0$ by showing that the right-hand side converges in probability to zero. We have $n^{-1}\mathbf{X}'\mathbf{Z} \xrightarrow{p} \mathbb{E}[\mathbf{x}_i\mathbf{z}_i']$, $n^{-1}\mathbf{Z}'\mathbf{Z} \xrightarrow{p} \mathbb{E}[\mathbf{z}_i\mathbf{z}_i']$, and $n^{-1}\mathbf{Z}'\boldsymbol{\varepsilon} \xrightarrow{p} \mathbb{E}[\mathbf{z}_i\boldsymbol{\varepsilon}_i] = \mathbf{0}$, where the last equality follows from the exogeneity condition of the instruments. Therefore, by Slutsky's theorem, $\hat{\beta}_{2\text{SLS}} \xrightarrow{p} \beta$. $\blacksquare$

Remark 3.2. As Wooldridge [16], Hayashi [11], and Hansen [10] emphasize, the consistency of the 2SLS estimator relies on the same conditions as the IV estimator, including the validity of the instruments. If the instruments are weakly correlated with the endogenous regressor, the 2SLS estimator will be less precise and may even be inconsistent. If the instruments are correlated with the error term, the 2SLS estimator will be biased and inconsistent. Therefore, it is essential to carefully select and validate instruments in applied work when using the 2SLS estimator.

Remark 3.3. As Amemiya [1], Hayashi [11], and Wooldridge [16] note, in the just-identified case, where the number of instruments equals the number of endogenous regressors, the 2SLS estimator is equivalent to the IV estimator. However, in the overidentified case, where there are more instruments than endogenous regressors, the 2SLS estimator provides a way to combine all the available instruments to achieve greater efficiency.

3.2.2 Asymptotic Normality

The asymptotic normality of the OLS estimator can be derived from the following theorem; see Hayashi [11], Wooldridge [16], and Hansen [10].

Theorem 3.4 (Asymptotic Normality of the OLS Estimator). *Under the assumptions of the linear model, random sampling, full rank, strict exogeneity, and homoskedasticity, the OLS estimator $\hat{\beta}_{\text{OLS}}$ is asymptotically normal, i.e.,*

$$\sqrt{n}(\hat{\beta}_{\text{OLS}} - \beta) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \sigma^2(\mathbb{E}[\mathbf{x}_i\mathbf{x}_i'])^{-1}), \quad (9)$$

where σ^2 is the variance of the error term ε_i .

As Amemiya [1], Hayashi [11], and Wooldridge [16] show, the asymptotic normality of the 2SLS estimator can be established under conditions analogous to those for OLS, together with instrument validity and sufficient strength. The asymptotic distribution of the 2SLS estimator is given by:

$$\sqrt{n}(\hat{\beta}_{2\text{SLS}} - \beta) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma}_{2\text{SLS}}), \quad (10)$$

where $\boldsymbol{\Sigma}_{2\text{SLS}}$ is the asymptotic variance-covariance matrix of the 2SLS estimator, which can be expressed in terms of the variance of the error term and the covariance between the instruments and the regressors. The exact form of $\boldsymbol{\Sigma}_{2\text{SLS}}$ can be derived using the delta method and the properties of the projection matrix $\mathbf{P}_Z$.

As Hayashi [11], Hansen [10], and Wooldridge [16] explain, asymptotic normality allows one to construct confidence intervals and perform hypothesis tests for the parameters of interest, using the appropriate IV/2SLS variance formula rather than the OLS one.

Remark 3.5. The asymptotic normality of the 2SLS estimator relies on the validity and strength of the instruments. If the instruments are weak, the asymptotic distribution of the 2SLS estimator may not be well-approximated by a normal distribution, and inference based on the normal approximation may be unreliable. In such cases, alternative methods for inference, such as the Anderson-Rubin test or the Stock-Yogo weak instrument test, may be more appropriate. Therefore, it is important to assess the strength of the instruments and consider alternative inference methods when using the 2SLS estimator in practice.

3.2.3 Efficiency

As Amemiya [1], Hayashi [11], and Wooldridge [16] show, the 2SLS estimator is more efficient than single-equation IV when multiple instruments are available because it combines all admissible instrument information. Under homoskedasticity, 2SLS is efficient within the class of IV estimators based on the same instrument set, in the sense of having weakly smaller asymptotic variance.

The efficiency gain from using the 2SLS estimator over the IV estimator can be substantial, especially when there are multiple instruments that are strongly correlated with the endogenous regressor. However, it is important to note that the efficiency of the 2SLS estimator depends on the validity and strength of the instruments, and if the instruments are weak or invalid, the 2SLS estimator may not be more efficient than the IV estimator, and may even be inconsistent. Amemiya [1], Hayashi [11], and Wooldridge [16] therefore emphasize careful consideration of the choice of instruments in practice.

3.3 Testing for Endogeneity

Before applying the 2SLS estimator, it is important to test for the presence of endogeneity in the model, since if the regressors are exogenous, the OLS estimator is consistent and more efficient than the 2SLS estimator, and therefore, there is no justification for using the 2SLS estimator. As Hansen [10] and Wooldridge [16] note, one common test for endogeneity is the Hausman test.

As Hansen [10] and Wooldridge [16] explain, this class of tests compares the OLS and 2SLS estimators to determine whether there is a systematic difference between them. Under the null hypothesis of exogeneity, both estimators are consistent, but OLS is efficient; under the alternative, OLS is inconsistent while IV/2SLS remains consistent if the instruments are valid. A standard Hausman statistic is

$$H = (\hat{\beta}_{OLS} - \hat{\beta}_{2SLS})^T (\text{Var}(\hat{\beta}_{OLS}) - \text{Var}(\hat{\beta}_{2SLS}))^{-1} (\hat{\beta}_{OLS} - \hat{\beta}_{2SLS}),$$

where $\text{Var}(\hat{\beta}_{OLS})$ and $\text{Var}(\hat{\beta}_{2SLS})$ are the asymptotic variances of the OLS and 2SLS estimators, respectively. Under standard regularity conditions, this statistic is asymptotically chi-squared with degrees of freedom equal to the number of tested coefficients. In practice, robust endogeneity diagnostics are often implemented as Durbin–Wu–Hausman regression-based tests, which remain valid under heteroskedasticity when paired with robust covariance estimation.

As Hansen [10] and Wooldridge [16] caution, the Hausman and Wu-Hausman tests are not definitive tests for endogeneity, because they rely on the maintained assumption that OLS is consistent under

the null. If that assumption fails, the test may not be valid. These diagnostics are therefore best used alongside broader specification and instrument-validity checks.

3.4 Weak Instruments

Weak instruments arise when the instruments are only weakly correlated with the endogenous regressors. This is not itself a form of endogeneity; rather, it is a failure of instrument relevance. Weak instruments can make 2SLS estimates imprecise and can also induce substantial finite-sample bias, causing IV estimates to drift toward OLS. As Hansen [10] and Chaussé [5] discuss, Stock and Yogo (2005) developed critical values based on first-stage diagnostics to assess instrument strength.

As Hansen [10] and Chaussé [5] note, the Stock-Yogo test provides critical values for the first-stage F -statistic that can be used to assess instrument strength. If the statistic is below the critical value, the instruments are judged weak; in that case, alternatives such as LIML or Anderson-Rubin-style inference may be preferable.

3.5 Overidentification and the Sargan-Hansen Test

When there are more instruments than endogenous regressors, the model is said to be overidentified. As Hayashi [11], Hall [8], and Hansen [10] explain, in this case it is possible to test the validity of the instruments using the Sargan-Hansen test for overidentifying restrictions. The Sargan-Hansen test statistic is given by:

$$J = n g_n(\hat{\beta})' \hat{\mathbf{W}} g_n(\hat{\beta}),$$

where $g_n(\hat{\beta}) = n^{-1} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\beta})$ is the vector of sample moment conditions evaluated at the estimator and $\hat{\mathbf{W}}$ is the weighting matrix used in estimation or testing. In the homoskedastic linear IV case, this reduces to the familiar Sargan form based on the IV residuals.

As Hayashi [11], Hall [8], and Hansen [10] show, under the null hypothesis that all overidentifying restrictions are valid, the Sargan-Hansen test statistic is asymptotically chi-squared with degrees of freedom equal to the number of overidentifying restrictions, that is, the number of moment conditions minus the number of estimated parameters. If the test statistic exceeds the relevant critical value, the null is rejected, suggesting that at least one instrument or maintained moment restriction is invalid.

4. Generalized Method of Moments (GMM) Estimator

The generalized method of moments (GMM) estimator generalizes the instrumental variables (IV) estimator by using moment conditions implied by the model's assumptions. Like IV and 2SLS, GMM relies on valid instruments and can accommodate multiple endogenous regressors, multiple instruments, and both just-identified and overidentified models. As Hayashi [11], Hall [8], and Hansen [10] explain, its broader appeal lies in its flexibility: GMM can be applied to linear and nonlinear models, allows for homoskedastic or heteroskedastic errors, and efficiently combines multiple moment conditions to estimate the parameters of interest.

Given the following linear regression model (in matrix form):

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad (11)$$

where $\mathbf{y}$ is the $n \times 1$ vector of the dependent variable, $\mathbf{X}$ is the $n \times k$ matrix of regressors (including the intercept, the endogenous regressor, and the exogenous regressors), $\boldsymbol{\beta}$ is the $k \times 1$ vector of parameters, and $\boldsymbol{\varepsilon}$ is the $n \times 1$ vector of error terms. Let $\mathbf{Z}$ be the $n \times m$ matrix of instruments, where $m \geq k$, and assume that the following moment conditions hold:

$$E[\mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})] = \mathbf{0}. \quad (12)$$

As Hayashi [11], Hall [8], and Hansen [10] define it, the GMM estimator $\hat{\boldsymbol{\beta}}_{\text{GMM}}$ is the value of $\boldsymbol{\beta}$ that minimizes the quadratic form of the sample moment conditions, weighted by a positive definite weighting matrix $\mathbf{W}_n$:

$$\hat{\boldsymbol{\beta}}_{\text{GMM}} = \arg \min_{\boldsymbol{\beta}} \left(\frac{1}{n} \sum_{i=1}^n \mathbf{z}_i(y_i - \mathbf{x}_i' \boldsymbol{\beta}) \right)' \mathbf{W}_n \left(\frac{1}{n} \sum_{i=1}^n \mathbf{z}_i(y_i - \mathbf{x}_i' \boldsymbol{\beta}) \right). \quad (13)$$

In matrix notation, the objective function can be written as:

$$Q_n(\boldsymbol{\beta}) = \left(\frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right)' \mathbf{W}_n \left(\frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right). \quad (14)$$

The GMM estimator is obtained by minimizing the objective function $Q_n(\boldsymbol{\beta})$ with respect to $\boldsymbol{\beta}$, and setting the derivative equal to zero as follows:

$$\frac{\partial Q_n(\boldsymbol{\beta})}{\partial \boldsymbol{\beta}} = -2 \left(\frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right)' \mathbf{W}_n \left(\frac{1}{n} \mathbf{Z}'\mathbf{X} \right) = \mathbf{0}. \quad (15)$$

From which the GMM estimator can be derived, step by step, as follows:

$$\begin{aligned} \left(\frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) \right)' \mathbf{W}_n \left(\frac{1}{n} \mathbf{Z}'\mathbf{X} \right) &= \mathbf{0} \\ \left(\frac{1}{n} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})' \mathbf{Z} \right) \mathbf{W}_n \left(\frac{1}{n} \mathbf{Z}'\mathbf{X} \right) &= \mathbf{0} \\ (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})' \mathbf{Z} \mathbf{W}_n (\mathbf{Z}'\mathbf{X}) &= \mathbf{0} \\ (\mathbf{y}' - \boldsymbol{\beta}'\mathbf{X}') \mathbf{Z} \mathbf{W}_n (\mathbf{Z}'\mathbf{X}) &= \mathbf{0} \\ \mathbf{y}' \mathbf{Z} \mathbf{W}_n (\mathbf{Z}'\mathbf{X}) - \boldsymbol{\beta}' \mathbf{X}' \mathbf{Z} \mathbf{W}_n (\mathbf{Z}'\mathbf{X}) &= \mathbf{0} \\ \boldsymbol{\beta}' \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X} &= \mathbf{y}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X} \\ (\boldsymbol{\beta}' \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X})' &= (\mathbf{y}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X})' \\ (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X})' \boldsymbol{\beta} &= (\mathbf{Z}'\mathbf{X})' \mathbf{W}_n' \mathbf{Z}'\mathbf{y} \\ \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X} \boldsymbol{\beta} &= \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{y} \\ \boldsymbol{\beta} &= (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}'\mathbf{y} \end{aligned}$$

In the second step, we used the fact that the transpose of a product of matrices is the product of their transposes in reverse order. In the third step, we multiplied both sides by n to eliminate the factor of $1/n$. In the fourth step, we rearranged the terms to isolate β on one side of the equation. In the fifth step, we took the transpose of both sides to facilitate the next step. In the sixth step, we used the fact that the transpose of a product of matrices is the product of their transposes in reverse order again. Finally, in the last step, we solved for β by pre-multiplying both sides by the inverse of $\mathbf{X}'\mathbf{Z}\mathbf{W}_n\mathbf{Z}'\mathbf{X}$.

From this derivation, we can see that the GMM estimator can be expressed as:

$$\hat{\beta}_{\text{GMM}} = (\mathbf{X}'\mathbf{Z}\mathbf{W}_n\mathbf{Z}'\mathbf{X})^{-1} \mathbf{X}'\mathbf{Z}\mathbf{W}_n\mathbf{Z}'\mathbf{y}.$$

4.1 The Weighting Matrix $\mathbf{W}_n$

As Hall [8], Hayashi [11], and Hansen [10] explain, the role of the weighting matrix $\mathbf{W}_n$ in the GMM objective function is to account for the variance-covariance structure of the moment conditions, weighting them inversely to their variance and accounting for potential correlation. In the first step of GMM estimation, any matrix that satisfies the following conditions can be used as the weighting matrix $\mathbf{W}_n$:

- (1) it must be positive definite, which ensures that the quadratic form objective function has a unique minimum;
- (2) it should converge in probability to a positive definite matrix $\mathbf{W}$ as the sample size increases, which ensures the consistency of the GMM estimator; and
- (3) its dimension must be $m \times m$ (where m is the number of moment conditions), which ensures that the objective function is properly defined.

In most applications, in the first step of GMM estimation, a simple choice for the weighting matrix is the identity matrix $\mathbf{W}_n = \mathbf{I}_m$, which gives equal weight to all moment conditions. While this choice yields a consistent estimator, it is generally not asymptotically efficient, as it does not account for the variance-covariance structure of the moment conditions (which may exhibit heteroskedasticity or correlation).

In practice, many matrices satisfy the conditions for $\mathbf{W}_n$, which implies there are many GMM estimators that are consistent for β . However, these estimators have different asymptotic efficiency properties. As Hayashi [11], Hall [8], and Hansen [10] show, the optimal weighting matrix is $\mathbf{W}_n = \mathbf{S}_n^{-1}$, where $\mathbf{S}_n$ is a consistent estimator of the variance-covariance matrix of the moment conditions. This optimal choice produces the asymptotically efficient GMM estimator, which achieves the lowest asymptotic variance among all GMM estimators based on the same set of moment conditions.

4.2 The Optimal Weighting Matrix

Definition 4.1 (Optimal Weighting Matrix). In the linear i.i.d. GMM setting, the optimal weighting matrix is the inverse of the asymptotic covariance matrix of the sample moments. A common heteroskedasticity-robust estimator takes the form $\mathbf{W}_n = \mathbf{S}_n^{-1}$, where $\mathbf{S}_n$ is a consistent estimator of that

covariance matrix, for example

$$\mathbf{S}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{z}_i \mathbf{z}_i^\top (y_i - \mathbf{x}_i^\top \hat{\boldsymbol{\beta}}_{\text{GMM}})^2.$$

With serial dependence, $\mathbf{S}_n$ should instead be replaced by a consistent long-run covariance estimator.

Proposition 4.2 (Efficiency of the Optimal GMM Estimator). *The GMM estimator that uses the optimal weighting matrix $\mathbf{W}_n = \mathbf{S}_n^{-1}$ is asymptotically efficient among the class of GMM estimators that use the same set of moment conditions. As Hayashi [11], Hall [8], and Hansen [10] show, this means that the asymptotic variance of the optimal GMM estimator is less than or equal to the asymptotic variance of any other GMM estimator that uses a different weighting matrix, in the positive semi-definite sense.*

Proof. Let $\mathbf{G}$ denote the Jacobian of the population moments and let Ω denote the asymptotic covariance matrix of $\sqrt{n}g_n(\beta_0)$. For a generic limiting weighting matrix $\mathbf{W}$, the asymptotic variance of the GMM estimator is

$$(\mathbf{G}'\mathbf{W}\mathbf{G})^{-1} \mathbf{G}'\mathbf{W}\Omega\mathbf{W}\mathbf{G}(\mathbf{G}'\mathbf{W}\mathbf{G})^{-1}.$$

When $\mathbf{W} = \Omega^{-1}$, this simplifies to

$$(\mathbf{G}'\Omega^{-1}\mathbf{G})^{-1},$$

which is the smallest variance in the positive semi-definite ordering among GMM estimators based on the same moment conditions. Hence the optimal-weight GMM estimator is asymptotically efficient within that class. \$■\$

Remark 4.3. The optimal GMM estimator is particularly useful in the presence of heteroskedasticity or autocorrelation in the error terms, as it accounts for the variance-covariance structure of the moment conditions, leading to more efficient estimates. However, it is important to note that the optimal GMM estimator relies on a consistent estimator of the variance-covariance matrix of the moment conditions, which can be challenging to obtain in practice, especially in the presence of heteroskedasticity or autocorrelation. In such cases, alternative methods for estimating the variance-covariance matrix, such as heteroskedasticity-robust or Newey–West estimators, may be necessary to ensure the efficiency of the GMM estimator.

4.3 Second Step GMM Estimation

In practice, the optimal GMM estimator is often implemented using a two-step procedure. As Hall [8], Hayashi [11], Hansen [10], and Chaussé [4] explain, the first step uses a consistent but potentially inefficient weighting matrix, such as the identity matrix $\mathbf{W}_n = \mathbf{I}_m$, to obtain an initial estimator and construct a consistent estimator of the variance-covariance matrix of the moment conditions, $\mathbf{S}_n$. In the second step, the optimal weighting matrix $\mathbf{W}_n = \mathbf{S}_n^{-1}$ is used to re-estimate the parameters, yielding the efficient GMM estimator. The two-step GMM estimator can be expressed as follows:

$$\hat{\beta}_{\text{GMM}}^{(2)} = (\mathbf{X}'\mathbf{Z}\mathbf{S}_n^{-1}\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}\mathbf{S}_n^{-1}\mathbf{Z}'\mathbf{y}, \quad (16)$$

where $\mathbf{S}_n$ is the consistent estimator of the variance-covariance matrix of the moment conditions obtained from the first step. The two-step GMM estimator is consistent and asymptotically efficient under the same conditions as the optimal GMM estimator, and it is widely used in practice due to its computational convenience and efficiency properties.

S is estimated using the residuals from the first step GMM estimation as follows:

$$\mathbf{S}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{z}_i \mathbf{z}_i' (y_i - \mathbf{x}_i' \hat{\beta}_{\text{GMM}}^{(1)})^2,$$

where $\hat{\beta}_{\text{GMM}}^{(1)}$ is the first-step GMM estimator obtained using a simple weighting matrix, such as the identity matrix $\mathbf{W}_n = \mathbf{I}_m$. The second-step GMM estimator $\hat{\beta}_{\text{GMM}}^{(2)}$ is then obtained by using the optimal weighting matrix $\mathbf{W}_n = \mathbf{S}_n^{-1}$ in the GMM estimation, which yields the efficient GMM estimator. The optimal weighting matrix $\mathbf{W}_n$ can be computed using the formula:

$$\mathbf{W}_n = \mathbf{S}_n^{-1} = \left(\frac{1}{n} \sum_{i=1}^n \mathbf{z}_i \mathbf{z}_i' (y_i - \mathbf{x}_i' \hat{\beta}_{\text{GMM}}^{(1)})^2 \right)^{-1}.$$

The optimal weighting matrix $\mathbf{W}_n$ can be computed using the formula:

$$\mathbf{W}_n = \mathbf{S}_n^{-1} = \left(\frac{1}{n} \sum_{i=1}^n \mathbf{z}_i \mathbf{z}_i' (y_i - \mathbf{x}_i' \hat{\beta}_{\text{GMM}}^{(1)})^2 \right)^{-1}.$$

4.4 Properties of the GMM Estimators

4.4.1 Consistency

As Hayashi [11], Hall [8], and Hansen [10] explain, the GMM estimator is consistent under the following conditions:

1. The population moment condition is correctly specified:

$$E[\mathbf{z}_i u_i(\beta_0)] = \mathbf{0}, \quad u_i(\beta_0) = y_i - \mathbf{x}_i' \beta_0.$$

2. The sample moments satisfy a uniform law of large numbers, so that

$$g_n(\beta) = \frac{1}{n} \sum_{i=1}^n \mathbf{z}_i \{y_i - \mathbf{x}_i' \beta\}$$

converges uniformly in probability to

$$g(\beta) = E[\mathbf{z}_i \{y_i - \mathbf{x}_i' \beta\}].$$

3. The parameter space is compact, or the estimator lies in a compact set with probability approaching one.
4. The weighting matrix satisfies

$$\mathbf{W}_n \xrightarrow{p} \mathbf{W},$$

where $\mathbf{W}$ is symmetric positive definite.

5. The true parameter β_0 is identified, meaning that

$$g(\beta) = \mathbf{0} \quad \text{if and only if} \quad \beta = \beta_0.$$

In the linear IV/GMM model, this requires an appropriate rank condition, for example

$$\text{rank}\{E(\mathbf{z}_i \mathbf{x}_i')\} = k,$$

where k is the number of parameters.

Proposition 4.4 (Consistency of the Linear GMM Estimator). *Let*

$$g_n(\beta) = \frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\beta)$$

and define the GMM estimator by

$$\hat{\beta}_{\text{GMM}} = \arg \min_{\beta \in \mathcal{B}} Q_n(\beta), \quad Q_n(\beta) = g_n(\beta)' \mathbf{W}_n g_n(\beta).$$

Suppose that $g_n(\beta)$ converges uniformly in probability to $g(\beta)$, that $\mathbf{W}_n \xrightarrow{p} \mathbf{W}$, where $\mathbf{W}$ is positive definite, and that the population criterion

$$Q(\beta) = g(\beta)' \mathbf{W} g(\beta)$$

has a unique minimum at β_0 . Then

$$\hat{\beta}_{\text{GMM}} \xrightarrow{p} \beta_0.$$

Proof. By the uniform law of large numbers,

$$\sup_{\beta \in \mathcal{B}} \|g_n(\beta) - g(\beta)\| \xrightarrow{p} 0.$$

Since $\mathbf{W}_n \xrightarrow{p} \mathbf{W}$, it follows that

$$\sup_{\beta \in \mathcal{B}} |Q_n(\beta) - Q(\beta)| \xrightarrow{p} 0,$$

where

$$Q(\beta) = g(\beta)' \mathbf{W} g(\beta).$$

By correct specification,

$$g(\beta_0) = \mathbf{0},$$

so that

$$Q(\beta_0) = 0.$$

Because $\mathbf{W}$ is positive definite and the parameter is identified,

$$Q(\beta) > 0 \quad \text{for all } \beta \neq \beta_0.$$

Thus $Q(\beta)$ has a unique minimum at β_0 . Uniform convergence of Q_n to Q , together with uniqueness of the population minimizer, implies by the extremum estimator consistency theorem that

$$\hat{\beta}_{\text{GMM}} \xrightarrow{p} \beta_0.$$

■

As Hayashi [11], Hall [8], and Hansen [10] note, in the linear IV/GMM model one can also prove consistency directly from the closed-form estimator rather than only through the general extremum-estimator argument. This approach exploits the linear structure of the model and the probability limits of the sample moment matrices. The GMM estimator can be written as follows:

$$\begin{aligned} \hat{\beta}_{\text{GMM}} &= (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{y} \\ &= \beta_0 + (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \varepsilon. \end{aligned}$$

Since β_0 is a constant vector, we can write its convergence to zero as follows:

$$\hat{\beta}_{\text{GMM}} - \beta_0 = (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \varepsilon.$$

We can show that $\text{plim} (\hat{\beta}_{\text{GMM}} - \beta_0) = 0$ by showing that the right-hand side converges in probability to zero. We have $n^{-1} \mathbf{X}' \mathbf{Z} \xrightarrow{p} \mathbb{E}[\mathbf{x}_i \mathbf{z}_i']$, $\mathbf{W}_n \xrightarrow{p} \mathbf{W}$, and $n^{-1} \mathbf{Z}' \varepsilon \xrightarrow{p} \mathbb{E}[\mathbf{z}_i \varepsilon_i] = \mathbf{0}$, where the last equality follows from the exogeneity condition of the instruments. Therefore,

$$\begin{aligned} (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \varepsilon &= \left(\frac{1}{n} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \frac{1}{n} \mathbf{Z}' \mathbf{X} \right)^{-1} \left(\frac{1}{n} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \frac{1}{n} \mathbf{Z}' \varepsilon \right) \\ &\xrightarrow{p} (\mathbb{E}[\mathbf{x}_i \mathbf{z}_i'] \mathbf{W} \mathbb{E}[\mathbf{z}_i \varepsilon_i])^{-1} (\mathbb{E}[\mathbf{x}_i \mathbf{z}_i'] \mathbf{W} \mathbf{0}) \\ &= \mathbf{0}. \end{aligned}$$

Thus, we have shown that $\text{plim} \hat{\beta}_{\text{GMM}} - \beta_0 = 0$, which implies that $\hat{\beta}_{\text{GMM}}$ is consistent for β_0 .

4.4.2 Asymptotic Normality

As Hayashi [11], Hall [8], and Hansen [10] show, the GMM estimator is asymptotically normal under the following regularity conditions:

- (i) The GMM estimator $\hat{\beta}_{\text{GMM}}$ is consistent for the true parameter vector β_0 .
- (ii) The moment conditions are correctly specified:

$$E[\mathbf{z}_i u_i(\beta_0)] = \mathbf{0}.$$

- (iii) The sample moments satisfy a central limit theorem:

$$\sqrt{n}g_n(\beta_0) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Omega),$$

where

$$g_n(\beta) = \frac{1}{n} \sum_{i=1}^n \mathbf{z}_i u_i(\beta),$$

and Ω is the long-run variance matrix of the moment conditions.

- (iv) The weighting matrix satisfies

$$\mathbf{W}_n \xrightarrow{p} \mathbf{W},$$

where $\mathbf{W}$ is symmetric positive definite.

- (v) The population moment function $g(\beta) = E[g_i(\beta)]$ is continuously differentiable in a neighborhood of β_0 , with Jacobian

$$\mathbf{G} = \left. \frac{\partial g(\beta)}{\partial \beta'} \right|_{\beta=\beta_0}.$$

- (vi) The rank condition holds:

$$\text{rank}(\mathbf{G}) = k,$$

where k is the number of parameters.

Under these conditions,

$$\sqrt{n}(\hat{\beta}_{\text{GMM}} - \beta_0) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \Sigma_{\text{GMM}}),$$

where

$$\Sigma_{\text{GMM}} = (\mathbf{G}'\mathbf{W}\mathbf{G})^{-1}\mathbf{G}'\mathbf{W}\Omega\mathbf{W}\mathbf{G}(\mathbf{G}'\mathbf{W}\mathbf{G})^{-1}.$$

For the linear IV/GMM model,

$$u_i(\beta) = y_i - \mathbf{x}_i' \beta,$$

so that

$$\mathbf{G} = E \left[\frac{\partial}{\partial \beta'} \mathbf{z}_i (y_i - \mathbf{x}_i' \beta) \right] = -E(\mathbf{z}_i \mathbf{x}_i').$$

Since the sign disappears in the sandwich variance formula, the asymptotic variance can equivalently be written using

$$\mathbf{Q}_{ZX} = E(\mathbf{z}_i \mathbf{x}_i').$$

Thus,

$$\Sigma_{\text{GMM}} = (\mathbf{Q}_{XZ} \mathbf{W} \mathbf{Q}_{ZX})^{-1} \mathbf{Q}_{XZ} \mathbf{W} \Omega \mathbf{W} \mathbf{Q}_{ZX} (\mathbf{Q}_{XZ} \mathbf{W} \mathbf{Q}_{ZX})^{-1},$$

where

$$\mathbf{Q}_{XZ} = E(\mathbf{x}_i \mathbf{z}_i') \quad \text{and} \quad \mathbf{Q}_{ZX} = E(\mathbf{z}_i \mathbf{x}_i').$$

This asymptotic normality result provides the basis for large-sample inference in the linear IV/GMM model. It allows one to construct confidence intervals and conduct hypothesis tests for the parameters of interest using the limiting distribution of the GMM estimator. As Hayashi [11], Hall [8], and Hansen [10] explain, the choice of weighting matrix affects the asymptotic variance of the estimator. The efficient GMM estimator is obtained when the limiting weighting matrix is chosen as $\mathbf{W} = \Omega^{-1}$, where Ω is the asymptotic variance of the moment conditions. In practice, Ω is replaced by a consistent estimator, often constructed from the sample analog of the moment covariance matrix, with heteroskedasticity- or autocorrelation-robust versions used when appropriate.

4.4.3 Efficiency

As Hayashi [11], Hall [8], and Hansen [10] show, among GMM estimators based on the same set of moment conditions, the asymptotically efficient estimator is obtained by choosing the limiting weighting matrix to be the inverse of the asymptotic variance-covariance matrix of the sample moments. Thus, the efficient weighting matrix is $\mathbf{W} = \Omega^{-1}$, where Ω denotes the asymptotic variance of $\sqrt{n}g_n(\beta_0)$. Under this optimal weighting matrix, the general GMM asymptotic variance formula reduces to

$$\Sigma_{\text{GMM}} = (\mathbf{G}' \Omega^{-1} \mathbf{G})^{-1}.$$

This is the smallest asymptotic variance attainable within the class of GMM estimators using the same moment conditions.

4.5 Tests

4.5.1 Overidentifying restrictions

The GMM framework provides several specification tests, including tests of overidentifying restrictions, tests of parameter restrictions, and more general tests of model specification. As Hall [8], Hansen [10], and Hayashi [11] explain, one of the most widely used specification tests in GMM is the Hansen (J)-test, also commonly referred to as the Sargan–Hansen test in the linear IV context. The test examines whether the sample moment conditions are sufficiently close to zero when evaluated at the GMM estimator.

Let

$$g_n(\beta) = \frac{1}{n} \sum_{i=1}^n g_i(\beta)$$

denote the vector of sample moments, and let $\hat{\beta}_{\text{GMM}}$ be the GMM estimator. The J -statistic is given by

$$J = n g_n(\hat{\beta}_{\text{GMM}})' \mathbf{W}_n g_n(\hat{\beta}_{\text{GMM}}),$$

where $\mathbf{W}_n$ is the weighting matrix used in the final-stage GMM criterion.

Under the null hypothesis that the overidentifying restrictions are valid, and provided that $\mathbf{W}_n$ consistently estimates the optimal weighting matrix Ω^{-1} , the statistic has the limiting distribution

$$J \xrightarrow{d} \chi_{q-k}^2,$$

where q is the number of moment conditions and k is the number of estimated parameters. Thus, $q - k$ is the number of overidentifying restrictions. In the linear IV model, this is often interpreted as the number of instruments excluded from the structural equation in excess of the number of endogenous regressors, or more generally as the number of instruments minus the number of parameters estimated.

As Hall [8], Hansen [10], and Hayashi [11] explain, if the value of J exceeds the relevant critical value from the χ_{q-k}^2 distribution, the null hypothesis is rejected. Rejection suggests that the moment conditions are not jointly compatible with the data and the maintained model. This may indicate that one or more instruments are invalid, but it may also reflect other forms of model misspecification, such as an incorrect functional form, omitted variables, or violations of the maintained error structure. The Hansen J -test is therefore an important diagnostic tool in applied GMM estimation, especially in overidentified models.

4.6 Other Diagnostic Tests

As Hall [8], Hansen [10], Chaussé [5], and Chaussé [4] emphasize, besides the test of overidentifying restrictions, several diagnostic tests are important in IV/GMM estimation. These include tests of endogeneity, such as the Durbin–Wu–Hausman test; tests of instrument relevance and weak identification, such as the first-stage F -statistic, the Cragg–Donald statistic, and the Kleibergen–Paap rank tests; tests

of underidentification; and Wald tests of parameter restrictions. In dynamic panel GMM, additional diagnostics are essential, including Arellano–Bond tests for serial correlation and Difference-in-Hansen tests for subsets of instruments. Researchers should also examine heteroskedasticity, serial correlation, robustness across alternative instrument sets, and the possibility of instrument proliferation. Together, these tests help assess whether the model is identified, whether the instruments are valid and relevant, and whether inference based on the GMM estimator is reliable.

5. Synthesis and Summary

This handout has examined the econometric problem of endogeneity and the main estimation strategies used to obtain consistent estimates when ordinary least squares is no longer reliable. The starting point is the classical linear model, in which OLS is consistent and asymptotically normal when the regressors are exogenous, that is, when the regressors are uncorrelated with the structural disturbance. Once this condition fails, OLS converges to the wrong probability limit, and the resulting inconsistency cannot be solved simply by increasing the sample size. Endogeneity is therefore an identification problem rather than merely a sampling problem.

The handout has emphasized three major sources of endogeneity: omitted variables, simultaneity, and measurement error. In each case, the core issue is that one or more regressors contain variation that is systematically related to the error term. For example, omitted ability in a returns-to-schooling equation may make schooling endogenous; simultaneous determination of price and quantity may make price endogenous in demand or supply equations; and classical measurement error may mechanically induce correlation between the observed regressor and the composite error term. These examples show that endogeneity is not a technical nuisance but a central threat to causal interpretation in applied economics.

Instrumental variables provide a way to restore identification by replacing the invalid orthogonality condition between regressors and errors with valid moment conditions involving instruments. A valid instrument must be relevant, meaning that it is sufficiently correlated with the endogenous regressor, and exogenous, meaning that it is uncorrelated with the structural error term. In the just-identified case, the IV estimator is obtained by solving the sample analogue of the instrument-error orthogonality condition. In overidentified models, where there are more instruments than endogenous regressors or parameters to be identified, two-stage least squares provides a systematic way to combine the available instruments by projecting the endogenous regressors onto the instrument space and then estimating the structural equation using the fitted values.

The generalized method of moments extends these ideas by treating IV and 2SLS as special cases of estimation based on moment conditions. GMM estimates the parameter vector by minimizing a quadratic form in the sample moments, with the weighting matrix determining how different moment conditions are combined. Any positive definite weighting matrix satisfying suitable regularity conditions can yield a consistent estimator, but the efficient GMM estimator is obtained by weighting the moments by the inverse of their asymptotic variance-covariance matrix. The handout also stresses that valid inference requires

diagnostic testing, including tests for endogeneity, weak instruments, underidentification, overidentifying restrictions, parameter restrictions, heteroskedasticity, serial correlation, and, in dynamic panel settings, serial correlation and instrument proliferation. Together, these tools allow researchers to assess whether the model is identified, whether the instruments are valid and relevant, and whether the resulting IV, 2SLS, or GMM estimates support reliable causal inference.

Table~1 summarizes the relationship among the simple instrumental variables estimator, the two-stage least squares estimator, and the generalized method of moments estimator. The three estimators share the same basic objective: to obtain consistent estimates when one or more regressors are endogenous by replacing invalid regressor-error orthogonality conditions with valid instrument-based moment conditions. The simple IV estimator is most transparent in the just-identified case, 2SLS extends the IV logic to settings with multiple instruments by projecting endogenous regressors onto the instrument space, and GMM provides the most general framework by estimating parameters from moment conditions using a weighting matrix. The table also highlights that the estimators differ mainly in how they use instruments and moment conditions, how they combine available information, and the assumptions required for consistency and efficient inference.

Table 1: Summary of the simple IV, 2SLS, and GMM estimators

Estimator	Core idea	Main formula	Key conditions and interpretation
Simple IV estimator	Uses an external source of exogenous variation to isolate the part of the endogenous regressor that is uncorrelated with the structural error term. It is most transparent in the just-identified case, where the number of instruments equals the number of parameters or endogenous components to be identified.	$\hat{\beta}_{IV} = (\mathbf{Z}'\mathbf{X})^{-1}\mathbf{Z}'\mathbf{y}$ when $\mathbf{Z}'\mathbf{X}$ is square and nonsingular.	Requires instrument relevance and exogeneity: $\text{rank}\{E(\mathbf{z}_i\mathbf{x}_i')\} = k$ and $E(\mathbf{z}_i\varepsilon_i) = \mathbf{0}$. The estimator is consistent but generally not unbiased in finite samples.
Two-stage least squares (2SLS)	Generalizes IV to settings with multiple instruments. The first stage projects the endogenous regressors onto the instrument space; the second stage estimates the structural equation using the projected regressors.	$\hat{\beta}_{2SLS} = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y}$, where $\mathbf{P}_Z = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'$.	Requires valid and sufficiently strong instruments. In the just-identified case, 2SLS coincides with the simple IV estimator. In the overidentified case, it combines all instruments and is efficient among IV estimators under homoskedasticity.
GMM estimator	Uses the sample analogue of population moment conditions and chooses the parameter vector that makes those sample moments as close to zero as possible, according to a weighting matrix. It nests IV and 2SLS as special cases.	$\hat{\beta}_{GMM} = \arg \min_{\beta} g_n(\beta)' \mathbf{W}_n g_n(\beta)$, where $g_n(\beta) = \frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\beta)$. For the linear model, $\hat{\beta}_{GMM} = (\mathbf{X}'\mathbf{Z}\mathbf{W}_n\mathbf{Z}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{Z}\mathbf{W}_n\mathbf{Z}'\mathbf{y}$.	Requires correctly specified moment conditions, identification, convergence of the sample moments, and a positive definite limiting weighting matrix. The efficient GMM estimator uses $\mathbf{W} = \mathbf{\Omega}^{-1}$, where $\mathbf{\Omega}$ is the asymptotic variance-covariance matrix of the moment conditions.

6. Review Questions and Exercises

Exercise 6.1 (Conceptual foundations). Explain why endogeneity is an identification problem rather than merely a small-sample problem. In your answer, distinguish between sampling variability and asymptotic bias, and explain why increasing the sample size does not generally solve the problem of endogenous regressors.

Exercise 6.2 (OLS under exogeneity). Consider the linear model

$$y_i = \mathbf{x}_i' \boldsymbol{\beta}_0 + \varepsilon_i.$$

Suppose that

$$\mathbb{E}[\mathbf{x}_i \varepsilon_i] = \mathbf{0} \quad \text{and} \quad \mathbb{E}[\mathbf{x}_i \mathbf{x}_i']$$

is finite and nonsingular.

(i) Show that

$$\hat{\boldsymbol{\beta}}_{\text{OLS}} = \boldsymbol{\beta}_0 + \left(\frac{1}{n} \mathbf{X}' \mathbf{X} \right)^{-1} \left(\frac{1}{n} \mathbf{X}' \boldsymbol{\varepsilon} \right).$$

(ii) Use the law of large numbers and Slutsky's theorem to show that

$$\hat{\boldsymbol{\beta}}_{\text{OLS}} \xrightarrow{p} \boldsymbol{\beta}_0.$$

Exercise 6.3 (OLS inconsistency under endogeneity). Suppose the linear model is correctly specified, but

$$\mathbb{E}[\mathbf{x}_i \varepsilon_i] = \boldsymbol{\delta} \neq \mathbf{0}.$$

Show that

$$\text{plim } \hat{\boldsymbol{\beta}}_{\text{OLS}} = \boldsymbol{\beta}_0 + \{\mathbb{E}(\mathbf{x}_i \mathbf{x}_i')\}^{-1} \boldsymbol{\delta}.$$

Explain why the second term represents asymptotic bias.

Exercise 6.4 (Sources of endogeneity). For each of the following cases, explain why the regressor of interest may be endogenous:

- (i) estimating the return to schooling when ability is unobserved;
- (ii) estimating demand when price and quantity are jointly determined;
- (iii) estimating the effect of income on consumption when income is measured with error;
- (iv) estimating the effect of class size on learning outcomes when parents sort into schools.

For each case, suggest one possible instrument and discuss whether it is likely to satisfy relevance and exogeneity.

Exercise 6.5 (Omitted variable bias). Suppose the true model is

$$y_i = \beta_0 + \beta_1 x_i + \gamma c_i + u_i, \quad \mathbb{E}[u_i | x_i, c_i] = 0,$$

but the researcher estimates

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i.$$

- (i) Write the composite error term ε_i .
- (ii) Derive $\text{Cov}(x_i, \varepsilon_i)$.
- (iii) Under what conditions is OLS inconsistent?
- (iv) Explain how the sign of the omitted variable bias depends on γ and $\text{Cov}(x_i, c_i)$.

Exercise 6.6 (Measurement error). Let the true model be

$$y_i = \beta_0 + \beta_1 x_i^* + u_i,$$

but suppose that the researcher observes

$$x_i = x_i^* + v_i,$$

where v_i is classical measurement error satisfying

$$\mathbb{E}[v_i] = 0, \quad \text{Cov}(x_i^*, v_i) = 0, \quad \text{Cov}(u_i, v_i) = 0.$$

- (i) Rewrite the model in terms of the observed regressor x_i .
- (ii) Show that the composite error is correlated with x_i .
- (iii) Derive the probability limit of the OLS estimator of β_1 .
- (iv) Explain why classical measurement error produces attenuation bias.

Exercise 6.7 (Definition of an instrumental variable). Let z_i be a candidate instrument for an endogenous regressor x_i in the model

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i.$$

- (i) State the relevance condition.
- (ii) State the exogeneity condition.
- (iii) Explain why relevance alone is not sufficient for consistency.
- (iv) Explain why exogeneity alone is not sufficient for identification.

Exercise 6.8 (Simple IV estimator). Consider the just-identified model

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i,$$

with instrument z_i , where

$$\mathbb{E}[z_i \varepsilon_i] = 0, \quad \text{Cov}(z_i, x_i) \neq 0.$$

Assume the variables have been demeaned so that no intercept is needed.

(i) Derive the population moment condition used to identify β_1 .

(ii) Show that

$$\beta_1 = \frac{\mathbb{E}[z_i y_i]}{\mathbb{E}[z_i x_i]}.$$

(iii) Derive the corresponding sample IV estimator.

(iv) Explain why the estimator is inconsistent if $\mathbb{E}[z_i \varepsilon_i] \neq 0$.

Exercise 6.9 (Matrix IV estimator). Consider the linear model

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon},$$

with instrument matrix $\mathbf{Z}$, where the model is just identified so that $\mathbf{Z}'\mathbf{X}$ is square and nonsingular.

(i) Starting from the sample moment condition

$$\frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\hat{\boldsymbol{\beta}}_{\text{IV}}) = \mathbf{0},$$

derive

$$\hat{\boldsymbol{\beta}}_{\text{IV}} = (\mathbf{Z}'\mathbf{X})^{-1} \mathbf{Z}'\mathbf{y}.$$

(ii) Show that

$$\hat{\boldsymbol{\beta}}_{\text{IV}} - \boldsymbol{\beta}_0 = (\mathbf{Z}'\mathbf{X})^{-1} \mathbf{Z}'\boldsymbol{\varepsilon}.$$

(iii) State the rank and exogeneity conditions needed for consistency.

Exercise 6.10 (First-stage regression and instrument strength). Suppose x_i is endogenous and z_i is proposed as an instrument. The first-stage equation is

$$x_i = \pi_0 + \pi_1 z_i + \mathbf{w}_i' \boldsymbol{\pi}_2 + v_i,$$

where $\mathbf{w}_i$ is a vector of included exogenous controls.

(i) What does the coefficient π_1 measure?

(ii) Why is testing $\pi_1 = 0$ important?

- (iii) Explain the meaning of a weak instrument.
- (iv) Discuss the consequences of weak instruments for bias, variance, and inference.

Exercise 6.11 (Deriving 2SLS). Let $\mathbf{Z}$ be the matrix of instruments and define the projection matrix

$$\mathbf{P}_Z = \mathbf{Z}(\mathbf{Z}'\mathbf{Z})^{-1}\mathbf{Z}'.$$

The 2SLS estimator is

$$\hat{\beta}_{2SLS} = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y}.$$

- (i) Explain the meaning of the projection matrix $\mathbf{P}_Z$.
- (ii) Show how $\mathbf{P}_Z\mathbf{X}$ represents the fitted values from projecting the regressors onto the instrument space.
- (iii) Explain why replacing endogenous regressors by their fitted values helps restore consistency.
- (iv) Under what condition does 2SLS reduce to the simple IV estimator?

Exercise 6.12 (Consistency of 2SLS). Using

$$\hat{\beta}_{2SLS} = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\mathbf{y},$$

and

$$\mathbf{y} = \mathbf{X}\beta_0 + \varepsilon,$$

show that

$$\hat{\beta}_{2SLS} - \beta_0 = (\mathbf{X}'\mathbf{P}_Z\mathbf{X})^{-1}\mathbf{X}'\mathbf{P}_Z\varepsilon.$$

Then express this result in terms of $\mathbf{Z}$ and show the probability limit of the right-hand side under instrument exogeneity and the appropriate rank condition.

Exercise 6.13 (GMM moment conditions). Consider the linear IV/GMM moment condition

$$\mathbb{E}[\mathbf{z}_i(y_i - \mathbf{x}_i'\beta_0)] = \mathbf{0}.$$

- (i) Define the sample moment function $g_n(\beta)$.
- (ii) Define the population moment function $g(\beta)$.
- (iii) Explain the role of the weighting matrix $\mathbf{W}_n$ in the GMM criterion.
- (iv) Explain why different positive definite weighting matrices can produce different GMM estimators.
- (v) Explain why all such estimators may be consistent but not equally efficient.

Exercise 6.14 (Closed-form linear GMM estimator). Let

$$g_n(\boldsymbol{\beta}) = \frac{1}{n} \mathbf{Z}'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

and define

$$Q_n(\boldsymbol{\beta}) = g_n(\boldsymbol{\beta})' \mathbf{W}_n g_n(\boldsymbol{\beta}),$$

where $\mathbf{W}_n$ is symmetric positive definite.

- (i) Differentiate $Q_n(\boldsymbol{\beta})$ with respect to $\boldsymbol{\beta}$.
- (ii) Set the derivative equal to zero.
- (iii) Derive the closed-form linear GMM estimator:

$$\hat{\boldsymbol{\beta}}_{\text{GMM}} = (\mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{X})^{-1} \mathbf{X}' \mathbf{Z} \mathbf{W}_n \mathbf{Z}' \mathbf{y}.$$

- (iv) Explain how this estimator relates to 2SLS when

$$\mathbf{W}_n = (\mathbf{Z}' \mathbf{Z})^{-1}.$$

Exercise 6.15 (Optimal GMM weighting matrix). Let

$$\sqrt{n} g_n(\boldsymbol{\beta}_0) \xrightarrow{d} \mathcal{N}(\mathbf{0}, \boldsymbol{\Omega}),$$

where $\boldsymbol{\Omega}$ is nonsingular.

- (i) State the optimal GMM weighting matrix.
- (ii) Show that when $\mathbf{W} = \boldsymbol{\Omega}^{-1}$, the asymptotic variance of the GMM estimator becomes

$$\boldsymbol{\Sigma}_{\text{GMM}} = (\mathbf{G}' \boldsymbol{\Omega}^{-1} \mathbf{G})^{-1}.$$

- (iii) Explain in words why moments with larger variance should receive smaller weight.
- (iv) Explain why the optimal GMM estimator is efficient only relative to estimators using the same set of moment conditions.

Exercise 6.16 (Diagnostic testing in IV/GMM). For each of the following tests, state the null hypothesis, the alternative hypothesis, and the practical interpretation of rejection:

- (i) Durbin–Wu–Hausman test for endogeneity;
- (ii) first-stage F -test for instrument relevance;
- (iii) Kleibergen–Paap underidentification test;
- (iv) Stock–Yogo weak-instrument test;
- (v) Hansen J -test for overidentifying restrictions.

Explain why no single diagnostic test is sufficient to establish that an IV/GMM model is credible.

A. Monte Carlo Simulation: OLS, IV, 2SLS, and GMM under Endogeneity

A.1 Simulation Results

This appendix uses a Monte Carlo experiment to illustrate the finite-sample behaviour of four estimators discussed in the handout: ordinary least squares (OLS), the simple instrumental variables (IV) estimator, two-stage least squares (2SLS), and generalized method of moments (GMM). The purpose is to show how endogeneity affects the distribution of the OLS estimator and how valid instruments can restore consistency. The simulation focuses on the coefficient of the endogenous regressor, β_2 , because this is the parameter most directly affected by correlation between the regressor and the structural error term.

The data-generating process is based on the linear model

$$y_i = \beta_1 + \beta_2 x_{2i} + \beta_3 x_{3i} + \varepsilon_i,$$

where x_{2i} is endogenous and x_{3i} is exogenous. The instruments for x_{2i} are z_{1i} and z_{2i} , both of which are constructed to be correlated with x_{2i} but uncorrelated with ε_i . The simulation uses sample size $N = 100$ and $R = 10,000$ replications. In each replication, OLS is estimated without instruments, the simple IV estimator uses only z_1 , while 2SLS and GMM use both instruments z_1 and z_2 . The empirical distribution of $\hat{\beta}_2$ is then summarized using density plots, means, variances, and biases.




Figure 1: Empirical densities of the Monte Carlo distribution of $\hat{\beta}_2$ for OLS, simple IV, 2SLS, and GMM.

The density plots show the central implication of the simulation. Because x_2 is constructed to be correlated with the structural disturbance ε , the OLS estimator of β_2 is centered away from the true value, indicating asymptotic bias caused by endogeneity. By contrast, the simple IV estimator, 2SLS estimator, and GMM estimator are centered much closer to the true value because they use instruments that are correlated with x_2 but uncorrelated with the error term. The simple IV estimator uses only one instrument and is therefore typically more dispersed, while 2SLS and GMM use both instruments and tend to produce tighter empirical distributions. In this design, 2SLS and GMM are expected to behave similarly, although GMM uses a second-step weighting matrix that accounts for the estimated variance-covariance structure of the moment conditions.

Table 2: Mean, variance, and bias of the Monte Carlo distribution of $\hat{\beta}_2$.

Estimator	Mean	Variance	Bias
OLS	1.3033	0.0043	0.3033
Simple IV	0.9869	0.0251	-0.0131
2SLS	1.0025	0.0107	0.0025
GMM	1.0024	0.0108	0.0024

The table summarizes the same lesson numerically. The mean of the OLS estimates is expected to differ substantially from the true value $\beta_2 = 1$, producing a non-negligible bias. This reflects the fact that OLS uses the endogenous variation in x_2 , including the component correlated with the disturbance term. The simple IV estimator should display much smaller bias because it uses only exogenous variation from z_1 , although its variance may be larger because it discards the additional information contained in z_2 . The 2SLS and GMM estimators use both instruments and therefore generally combine low bias with improved precision. The relative performance of 2SLS and GMM depends on the weighting matrix and the variance structure of the moment conditions; in this homoskedastic design, their results are likely to be close, while under heteroskedasticity the efficient GMM estimator may offer clearer efficiency gains.

A.2 R Code

The following code reproduces the Monte Carlo simulation, the density figure, and the summary table reported above.

```
# Load the packages used for simulation, reshaping, plotting, and tables.
library(ggplot2)
library(dplyr)
library(tidyr)
library(knitr)

set.seed(12345)

# Monte Carlo design
N <- 100
R <- 10000

# True parameters
beta_1 <- 1.0
beta_2 <- 1.0
```

```

beta_3 <- 0.5

# Strength of endogeneity
rho <- 0.60

# Storage
results <- matrix(NA_real_, nrow = R, ncol = 4)
colnames(results) <- c("OLS", "Simple IV", "2SLS", "GMM")

# Define a simple two-step linear GMM estimator.
gmm_linear <- function(y, X, Z) {

  n <- nrow(X)

  # First-step GMM using 2SLS weighting matrix
  W1 <- solve(crossprod(Z) / n)
  beta_1step <- solve(
    crossprod(X, Z) %*% W1 %*% crossprod(Z, X),
    crossprod(X, Z) %*% W1 %*% crossprod(Z, y)
  )

  # Residuals from first step
  uhat <- as.vector(y - X %*% beta_1step)

  # Robust estimate of moment variance
  Zu <- Z * uhat
  S <- crossprod(Zu) / n

  # Small ridge adjustment for numerical stability
  S <- S + diag(1e-8, ncol(S))

  # Second-step optimal weighting matrix
  W2 <- solve(S)

  beta_2step <- solve(
    crossprod(X, Z) %*% W2 %*% crossprod(Z, X),
    crossprod(X, Z) %*% W2 %*% crossprod(Z, y)
  )
}

```

```

    as.vector(beta_2step)
}

for (r in 1:R) {

  # Instruments and exogenous regressor
  z1 <- rnorm(N)
  z2 <- rnorm(N)
  x3 <- rnorm(N)

  # Components generating endogeneity
  v <- rnorm(N)
  e <- rnorm(N)

  # Structural error correlated with v
  eps <- rho * v + sqrt(1 - rho^2) * e

  # Generate an endogenous regressor using instruments and the latent shock
  #  $\epsilon \leftarrow v$ .
  x2 <- 0.7 * z1 + 0.7 * z2 + 0.5 * x3 + v

  # Outcome variable
  y <- beta_1 + beta_2 * x2 + beta_3 * x3 + eps

  # Matrices
  X <- cbind(1, x2, x3)

  # OLS
  beta_ols <- solve(crossprod(X), crossprod(X, y))

  # Simple IV: just identified, using z1 plus included exogenous regressors.
  Z_iv <- cbind(1, z1, x3)
  beta_iv <- solve(crossprod(Z_iv, X), crossprod(Z_iv, y))

  # 2SLS: use both instruments plus included exogenous regressors.
  Z_all <- cbind(1, z1, z2, x3)
  Pz <- Z_all %*% solve(crossprod(Z_all)) %*% t(Z_all)
}

```

```

beta_2sls <- solve(
  crossprod(X, Pz %*% X),
  crossprod(X, Pz %*% y)
)

# Two-step GMM: use the same full instrument set.
beta_gmm <- gmm_linear(y = y, X = X, Z = Z_all)

# Store beta_2 estimates
results[r, ] <- c(beta_ols[2], beta_iv[2], beta_2sls[2], beta_gmm[2])
}

results_df <- as.data.frame(results)

# Reshape the simulation output into long form for plotting.
density_df <- results_df |>
  pivot_longer(
    cols = everything(),
    names_to = "Estimator",
    values_to = "beta2_hat"
  )

density_df$Estimator <- factor(
  density_df$Estimator,
  levels = c("OLS", "Simple IV", "2SLS", "GMM")
)

# Plot the empirical distribution of the target coefficient by estimator.
ggplot(density_df, aes(x = beta2_hat)) +
  geom_density(linewidth = 0.8) +
  geom_vline(xintercept = beta_2, linetype = "dashed") +
  facet_wrap(~ Estimator, nrow = 2, ncol = 2, scales = "free_y") +
  labs(
    x = expression(hat(beta)[2]),
    y = "Empirical density"
  ) +
  theme_classic()

```

```
# Compute Monte Carlo means, variances, and biases by estimator.
summary_table <- density_df |>
  group_by(Estimator) |>
  summarise(
    Mean = mean(beta2_hat),
    Variance = var(beta2_hat),
    Bias = mean(beta2_hat) - beta_2,
    .groups = "drop"
  )

summary_table |>
  mutate(
    Mean = round(Mean, 4),
    Variance = round(Variance, 4),
    Bias = round(Bias, 4)
  ) |>
  kable(
    booktabs = TRUE,
    align = c("l", "r", "r", "r")
  )
```

A.3 Documentation of the R Code

This subsection explains the code presented in Appendix A.2 in the same order in which it is written and executed. The code naturally falls into three parts: the simulation itself, the construction of the density figure, and the construction of the summary table.

The first code block sets up and runs the Monte Carlo experiment. It loads the packages required for plotting, reshaping data, and formatting tables; fixes the random seed for reproducibility; and defines the simulation dimensions $N = 100$ and $R = 10000$. It then specifies the true parameter values β_1 , β_2 , and β_3 , together with the endogeneity parameter ρ . Inside the simulation loop, the code generates two instruments z_1 and z_2 , one exogenous regressor x_3 , and latent shocks v and e . The structural disturbance ε is constructed from v and e , so that ε is correlated with the endogenous regressor x_2 . After generating the outcome variable y , the code estimates the coefficient on x_2 using OLS, a simple just-identified IV estimator, 2SLS, and a two-step GMM estimator. The estimate of β_2 from each method is stored in the results matrix for later comparison.

The second code block turns the simulation output into the figure reported in Appendix A.1. The matrix of simulation results is first converted into a data frame and then reshaped from wide form to long form using `pivot_longer()`. This produces one column for the estimator label and one column for the simulated values of $\hat{\beta}_2$. The estimator labels are then converted into an ordered factor so that the panels appear in a consistent order. Finally, `ggplot()` is used to draw empirical density curves for the four estimators, with a dashed vertical line marking the true value β_2 . The `facet_wrap()` call arranges the four densities in a 2×2 layout, making it easy to compare the sampling distributions visually.

The third code block constructs the summary table reported in Appendix A.1. It groups the long-form simulation data by estimator and computes three Monte Carlo summaries for $\hat{\beta}_2$: the mean, the variance, and the bias relative to the true coefficient β_2 . The bias is defined as the average estimate minus the true parameter, so it shows whether an estimator is centered on the target. After these quantities are computed, the code rounds them to four decimal places and passes the result to `kable()` with `booktabs` styling, while Quarto supplies the table caption and placement metadata through the chunk options.

Bibliography

- [1] Takeshi Amemiya. *Advanced Econometrics*. Cambridge, MA: Harvard University Press, 1985. ISBN: 9780674005600.
- [2] Colin A. Cameron and Pravin K. Trivedi. *Microeconometrics: Methods and Applications*. Cambridge, UK: Cambridge University Press, 2005. ISBN: 9780521848053.
- [3] David Card. *Using Geographic Variation in College Proximity to Estimate the Return to Schooling*. Princeton University, Industrial Relations Section Working Paper No. 317. Available at: https://davidcard.berkeley.edu/papers/geo_var_schooling.pdf. 1993. URL: https://davidcard.berkeley.edu/papers/geo%5C_var%5C_schooling.pdf.
- [4] Pierre Chaussé. *Generalized Method of Moments with R*. <https://cran.r-project.org/web/packages/momentfit/vignettes/gmmS4.pdf>. 2025.
- [5] Pierre Chaussé. *Instrumental Variables Tools for the Case of Weak and/or Many Instruments*. <https://cran.r-project.org/web/packages/momentfit/vignettes/weak.pdf>. 2025.
- [6] James Davidson. *Econometric Theory*. Oxford, UK: Wiley-Blackwell, 2000. ISBN: 9780631215844.
- [7] Russell Davidson and James G. MacKinnon. *Econometric Theory and Methods*. New York: Oxford University Press, 2004. ISBN: 9780195123722.
- [8] Alastair R. Hall. *Generalized Method of Moments*. Oxford, UK: Oxford University Press, 2005. ISBN: 9780198775201.
- [9] James D. Hamilton. *Time Series Analysis*. Princeton, NJ: Princeton University Press, 1994. ISBN: 9780691042893.
- [10] Bruce E. Hansen. *Econometrics*. Princeton, NJ: Princeton University Press, 2022. ISBN: 9780691235899.
- [11] Fumio Hayashi. *Econometrics*. Princeton, NJ: Princeton University Press, 2000. ISBN: 9780691010182.
- [12] George G. Judge et al. *Introduction to the Theory and Practice of Econometrics*. 2nd ed. New York: John Wiley & Sons, 1988. ISBN: 9780471624141.
- [13] M. Hashem Pesaran. *Time Series and Panel Data Econometrics*. Oxford, UK: Oxford University Press, 2015. ISBN: 9780198736912.
- [14] John Stachurski. *A Primer in Econometric Theory*. Cambridge, MA: The MIT Press, 2016. ISBN: 9780262034906.
- [15] Halbert White. *Asymptotic Theory for Econometricians*. Emerald Group/Academic Press, 2000. ISBN: 9780127466521.

- [16] Jeffrey M. Wooldridge. *Econometric Analysis of Cross Section and Panel Data*. 2nd ed. Cambridge, MA: The MIT Press, 2010. ISBN: 9780262232586.